

# Preserving Reciprocal Consistency in Distributed Graph Databases

---

Jack Waudby<sup>1</sup>, Paul Ezhilchelvan<sup>1</sup>, Jim Webber<sup>2</sup> & Isi Mitrani<sup>1</sup>

April 27, 2020

<sup>1</sup>Newcastle University

<sup>2</sup>Neo4j



# Graph Databases

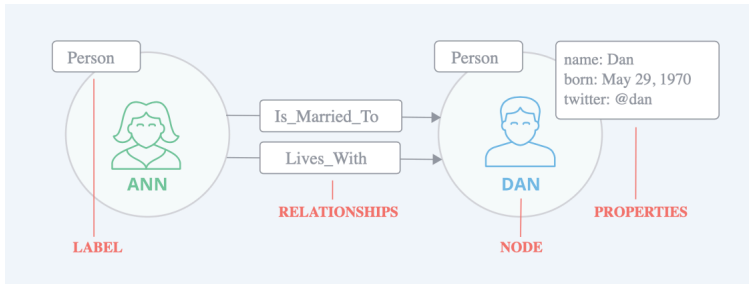


Figure 1: Labeled property graph<sup>1</sup>

---

<sup>1</sup><https://neo4j.com>

- **Efficient graph analysis:** reachability, pattern matching, shortest path search and clustering

- **Efficient graph analysis:** reachability, pattern matching, shortest path search and clustering
- **Use cases:** telecommunications, pharma, publishing, finance, social media

# Graph Database Popularity




Figure 2: Database popularity by data model<sup>2</sup>

<sup>1</sup>[https://db-engines.com/en/ranking\\_categories](https://db-engines.com/en/ranking_categories)

- Graph exceeds the storage capacity of a single machine

- Graph exceeds the storage capacity of a single machine
- Partition the graph across multiple machines in a cluster, with partitions replicated for fault-tolerance and availability

- Graph exceeds the storage capacity of a single machine
- Partition the graph across multiple machines in a cluster, with partitions replicated for fault-tolerance and availability
- Graph partitioning inevitably leads to **distributed edges**







Figure 3: Connected graph

# Distributed Graph Databases: Partitioning



**Figure 4:** Graph partitioned across 3 servers

## Reciprocal Consistency

A logical edge is represented by 2 physical records



Figure 5: A reciprocal consistent edge

# Reciprocal Consistency

A logical edge is represented by 2 physical records





Figure 5: A reciprocal consistent edge

# Reciprocal Consistency Violation

Two concurrent transactions:

- Mr Red requests to book the seat,  $T_R$
- Mr Blue requests to book the seat,  $T_B$

# Reciprocal Consistency Violation



**Figure 6:** Concurrent transactions interleave and violate reciprocal consistency

# Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

# Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency



# Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

# Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

# Reciprocal Consistency Violation




Figure 7: Reciprocal inconsistent distributed edge

# Reciprocal Consistency Violation




Figure 7: Reciprocal inconsistent distributed edge

# Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
  - A distributed edge's reciprocal information is separated by the network

---

<sup>3</sup>Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

# Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
  - A distributed edge's reciprocal information is separated by the network
- Reciprocal inconsistency is a source of corruption<sup>3</sup>

---

<sup>3</sup>Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

# Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
  - A distributed edge's reciprocal information is separated by the network
- Reciprocal inconsistency is a source of corruption<sup>3</sup>
- No known concurrency control protocol exists specific to graph databases and maintaining reciprocal consistency

---

<sup>3</sup>Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

## Our Solution: Collision Detection

- Carefully abort transaction(s)



## Our Solution: Collision Detection

- Carefully abort transaction(s)
- No centralized control or synchronized clock

## Our Solution: Collision Detection


- Carefully abort transaction(s)
- No centralized control or synchronized clock
- Permits interleavings that preserve reciprocal consistency

- Updates are initially provisional


- Updates are initially provisional
- Transactions distinguish between their first and second updates using labels

- Updates are initially provisional
- Transactions distinguish between their first and second updates using labels
- When a transaction attempts to update a given record, it can identify all other transactions that have earlier updated that record provisionally


# Interference-free Update



# Interference-free Update




# Interference-free Update






# Interference-free Update



Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort


# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort



# Collision Detection


Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort






# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort



# Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort



A second consistency problem related to distributed edges:



A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges

A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges
- Each update is reciprocally consistent

A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges
- Each update is reciprocally consistent
- Then if  $T_R$  updates before  $T_B$  on a given edge then this order should be preserved across all edges

# Edge-Order Consistency Violation

Consider two edges,




Figure 8: Two distributed edges

# Edge-Order Consistency Violation

Two concurrent transactions by a travel agent:

- Requests to book room and seat for Mr Red,  $T_R$
- Requests to book room and seat for Mr Blue,  $T_B$

# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation



# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

# Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

# Edge-Order Consistency Violation




Figure 10: Edge-order consistency violation

- Transactions collect predecessors from each update

## Our Solution: Order Arbitration

- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state

## Our Solution: Order Arbitration



- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state
- Global state used to detect edge-order violations

## Our Solution: Order Arbitration

- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state
- Global state used to detect edge-order violations
- Offending transactions are aborted




Rule: Collect predecessors after every successful update





# Order Arbitration

Rule: Collect predecessors after every successful update





# Order Arbitration

Rule: Collect predecessors after every successful update




# Order Arbitration

Rule: Collect predecessors after every successful update




Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list

*Arbiter*




Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



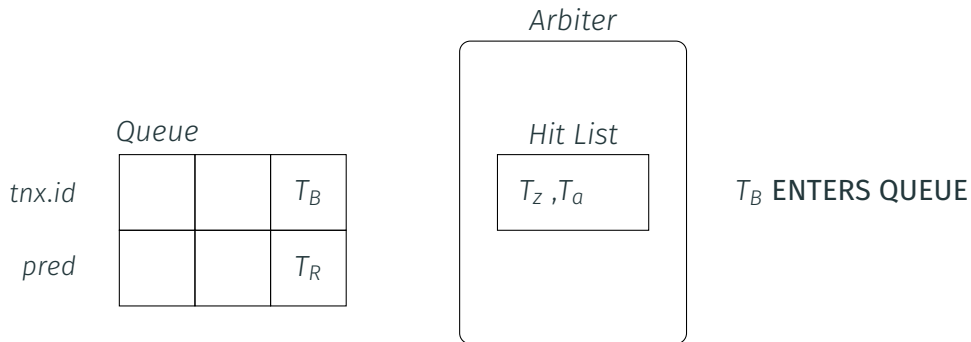
## Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



# Order Arbitration

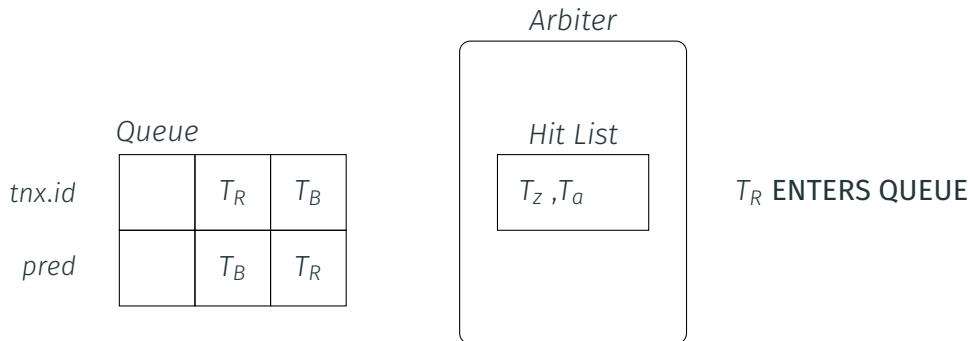
Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list





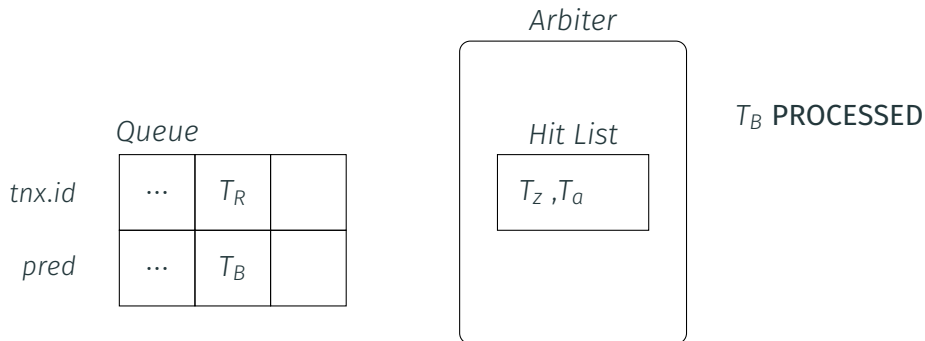
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



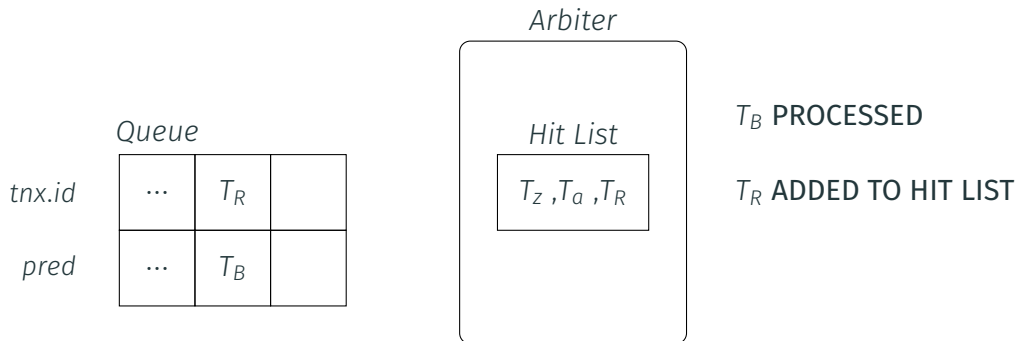
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



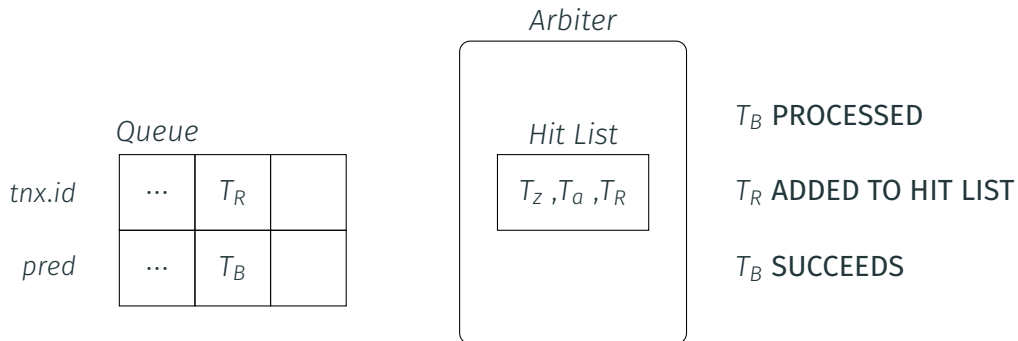
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



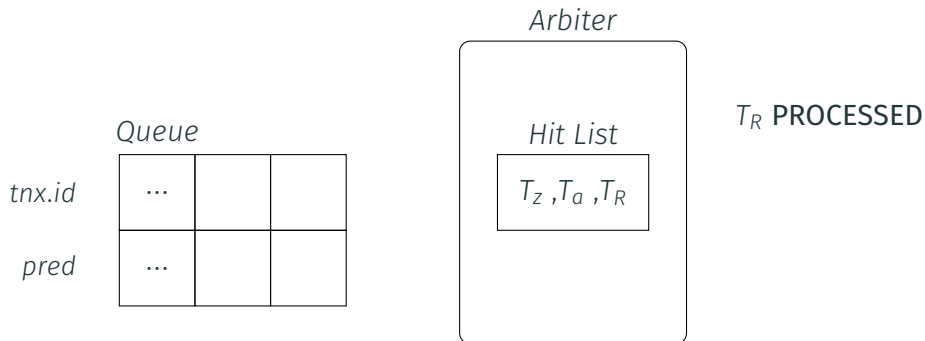
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



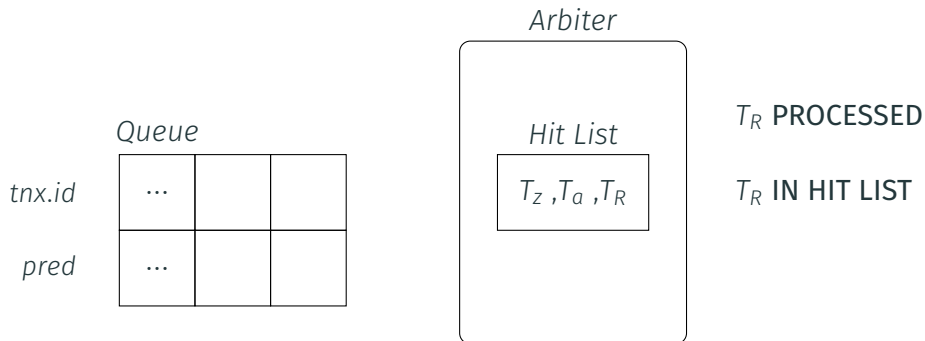
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



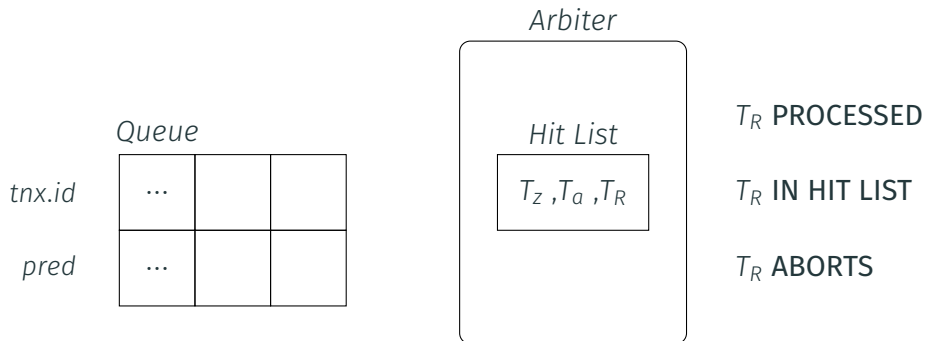
# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



# Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



If a transaction updates 1 or more distributed edges:

- Collect all predecessors



If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter

If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter
- If not in hit list then proceed and enter predecessors into the hit list

If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter
- If not in hit list then proceed and enter predecessors into the hit list
- Else in hit list then abort

1. Collision detection: enforces **reciprocal consistency** for distributed edges
2. Order arbitration: enforces **edge-order consistency** between transactions

Performance measures of interest:

- Average number of transactions that are aborted
- Load at the arbiter

1. Developed approximate model
2. Measure accuracy of model through simulation

- Database size
- Transaction arrival rate
- Average updates per transaction
- Average arbiter service time
- Average network delay

## Aborts per second ( $R$ ) vs Database Size ( $N$ )




Figure 11: TPS = 1000, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 5ms



## Aborts per second ( $R$ ) vs Transaction Arrival Rate ( $\lambda$ )

Arbiter queue unstable at  $\sim 1500$  TPS




Figure 12: size = 10K, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 5ms

## Aborts per second ( $R$ ) vs Transaction Arrival Rate ( $\lambda$ )

Arbiter queue unstable at  $\sim 1100$  TPS




Figure 13: size = 10K, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 10ms

## Aborts per second ( $R$ ) vs Transaction Arrival Rate ( $\lambda$ )

Arbiter queue unstable at  $\sim 550$  TPS




Figure 14: size = 10K, av updates/tnx = 10, av arbiter service time = 10ms, av network delay = 5ms

- Model accuracy similar to simulations under a variety of parameter settings

- Model accuracy similar to simulations under a variety of parameter settings
- Between 1-4% transactions are aborted

- Model accuracy similar to simulations under a variety of parameter settings
- Between 1-4% transactions are aborted
- Improve accuracy of simulation

Thanks for listening!

Any Questions?

Email: [j.waudby2@newcastle.ac.uk](mailto:j.waudby2@newcastle.ac.uk)

Twitter: [@waudberry\\_7](https://twitter.com/waudberry_7)