

Preserving Reciprocal Consistency in Distributed Graph Databases

Jack Waudby¹, Paul Ezhilchelvan¹, Jim Webber² & Isi Mitrani¹

April 27, 2020

¹Newcastle University

²Neo4j



Property Graph

- Graph databases model data as a *property graph*

Property Graph

- Graph databases model data as a *property graph*
- Vertices represent entities and edges the relationships between entities

Property Graph

- Graph databases model data as a *property graph*
- Vertices represent entities and edges the relationships between entities
- Edges are **always** directional

Property Graph

- Graph databases model data as a *property graph*
- Vertices represent entities and edges the relationships between entities
- Edges are **always** directional




Figure 1: Vertices connected by an edge

Storage Layer Representation

- In the storage layer,
 - edge directionality does **not** exist
 - connected vertices store information about each other

Storage Layer Representation

- In the storage layer,
 - edge directionality does **not** exist
 - connected vertices store information about each other
- Bi-directional edge traversal speeds up query performance

Storage Layer Representation

- In the storage layer,
 - edge directionality does **not** exist
 - connected vertices store information about each other
- Bi-directional edge traversal speeds up query performance




Figure 2: Edge storage layer representation

When the adjacency list edge pointers for a given edge refer to each other in a complementary manner, that edge is *reciprocally consistent*

Partition graph across machines in a cluster




Figure 3: Partitioned graph

A degree of concurrency control is needed for ensuring reciprocal consistency of distributed edges

Reciprocal Inconsistency

Distributed edge ab indicates Tolkien wrote The Hobbit




Figure 4: Distributed edge

Reciprocal Inconsistency

S_i




S_j




Reciprocal Inconsistency

- T_x deletes the edge



Reciprocal Inconsistency

- T_x deletes the edge
- T_y appends a property `{year:1937}`



Reciprocal Inconsistency

The distributed edge is now *reciprocally inconsistent*

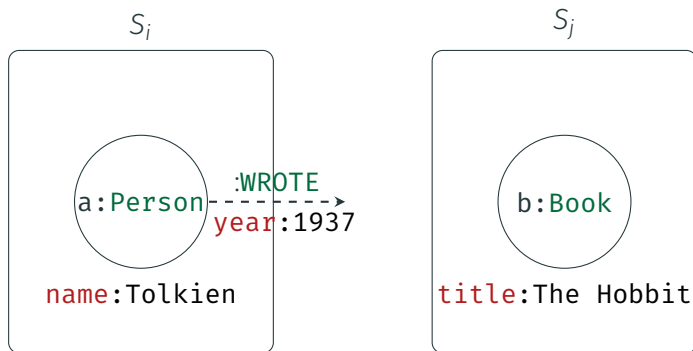


Figure 5: Reciprocally inconsistent distributed edge

Reciprocal Inconsistency

Storage layer consists of two inconsistent uni-directional edge pointers



(a) S_i



(b) S_j

Figure 6: Edge storage layer representation

- Reciprocally inconsistent edges are the source for *semantic corruption*

- Reciprocally inconsistent edges are the source for *semantic corruption*
- Semantic corruption spreads through the database

Reciprocal Inconsistency

- Reciprocally inconsistent edges are the source for *semantic corruption*
- Semantic corruption spreads through the database
- Motivated the design of a lightweight concurrency control protocol that preserves distributed edge reciprocal consistency

Design considerations:

Protocol Design Considerations

Design considerations:

1. Graph workloads exhibit high contention

Design considerations:

1. Graph workloads exhibit high contention
2. Graph transactions tend to be longer-lived than those in other databases

Design considerations:

1. Graph workloads exhibit high contention
2. Graph transactions tend to be longer-lived than those in other databases

Protocol must permit multiple updates on the same distributed edge provided they are *sufficiently* apart in time to ensure reciprocal consistency

- **Fact:** a transaction updating a distributed edge must update one edge pointer then **immediately** update the other

- **Fact:** a transaction updating a distributed edge must update one edge pointer then **immediately** update the other
- **Rule:** an update is permitted if the immediately preceding update was done at least Δ time before. Else, abort

- **Fact:** a transaction updating a distributed edge must update one edge pointer then **immediately** update the other
- **Rule:** an update is permitted if the immediately preceding update was done at least Δ time before. Else, abort
- **Assumption:** the time interval that elapses between completing an update at one edge pointer and starting at the other can be estimated to be δ .

- **Fact:** a transaction updating a distributed edge must update one edge pointer then **immediately** update the other
- **Rule:** an update is permitted if the immediately preceding update was done at least Δ time before. Else, abort
- **Assumption:** the time interval that elapses between completing an update at one edge pointer and starting at the other can be estimated to be δ .
- **Parameter Selection:** Choose $\Delta > \delta$

Delta Protocol: Example


Rule: an update is permitted if the preceding update was done at least Δ time before.

Else, abort

Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.


Else, abort



Delta Protocol: Example

Rule: an update is permitted if the preceding update was done at least Δ time before.

Else, abort



- Reciprocal consistency is preserved if the time taken to complete an update at one edge pointer and start at other remains less than Δ

- Reciprocal consistency is preserved if the time taken to complete an update at one edge pointer and start at other remains less than Δ
- If Δ is exceeded then reciprocal inconsistency and hence semantic corruption can occur


- Reciprocal consistency is preserved if the time taken to complete an update at one edge pointer and start at other remains less than Δ
- If Δ is exceeded then reciprocal inconsistency and hence semantic corruption can occur
- Setting a large Δ tends to preserve consistency but leads to more aborted transactions

Through simulations the following two metrics for various values of Δ were evaluated:

- Time taken for 10% of a large database to become semantically corrupt
- Fraction of transactions aborted per second


Time until 10% semantic corruption $\log(U)$ vs Transaction Arrival Rate (λ)

For $\Delta = 50ms$, time taken for 10% of a large database to become semantically corrupt is between to 1-75 years



Fraction of Aborts (α) vs Transaction Arrival Rate (λ)

For $\Delta = 50ms$, the fraction of aborts is between 1 – 5%



- Lack of concurrency control in a distributed graph database can lead to reciprocally inconsistent distributed edges
- Resulting in the spread of semantic corruption
- Delta protocol prevents reciprocal inconsistency given Δ is not exceeded
- Delta protocol significantly reduces the spread of semantic corruption at the cost of a very small fraction of aborts